

# Parallel Programming Lab

## 01 - Introduction to CUDA

---

Filippo Ziche (filippo.ziche@univr.it)

University of Verona - Department of Engineering for Innovation Medicine  
PARCO - Parallel Computing Lab

March 4, 2025

# Agenda

- What is CUDA?
- Calling a Device Functions
- CUDA Concepts and Keywords
- Memory Management
- Error handling
- CUDA Compiler
- Exercises!

# What is CUDA?

- CUDA (**C**ompute **U**nified **D**evice **A**rchitecture) is a parallel computing platform and programming model that allows to use a GPU for general purpose computing.
- It is small extension of C/C++ language.
- It allows you to accelerate your C/C++ code by moving the computationally intensive portions of your code to an NVIDIA GPU.

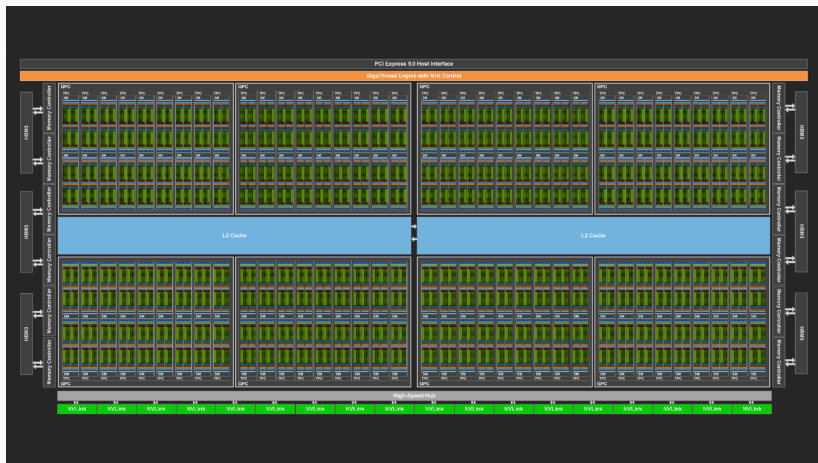
# Calling a Device Function (Kernel)

- A **kernel** is a function callable from the host and executed on the CUDA device. They are the heart of your CUDA code.

```
__global__ void my_kernel(int* x, int N) { ... }  
int main() {  
    ...  
    my_kernel<<<blocks, threads>>>(x, N);  
    ...  
}
```

- A kernel is defined using the **\_\_global\_\_** declaration specifier. They must have return type **void**.
- The number of CUDA threads that execute a kernel is specified using a new **<<< ... >>>** execution configuration syntax

# CUDA Architecture Overview



A lot of concurrent threads.

# Built-in Variables

CUDA provides a set of built-in variables to facilitate the **mapping between threads and data**. In the case of 1-dimensional data we have the following:

- **threadIdx.x** Thread index within the block
- **blockIdx.x** Block index within the grid
- **blockDim.x** Dimension of the block (number of threads)
- **gridDim.x** Dimension of the grid (number of blocks)

Common pattern to get **global id** of the thread:

```
gid = blockIdx.x * blockDim.x + threadIdx.x;
```

# Common CUDA Code Organization

1. Initialize host data structures
2. Allocate device data structures (`cudaMalloc`)
3. Copy host data structures to device (`cudaMemcpy`)
4. Invoke kernel function
5. Copy device result to host (`cudaMemcpy`)
6. Free host and device data structures (`freeCuda`)

# CUDA Memory Management I

- Allocate memory on the device.

```
__host__ __device__ cudaError_t cudaMalloc(void** devPtr, size_t size)
```

- Copies data between host and device.

```
__host__ cudaError_t cudaMemcpy(void* dst, const void* src, size_t count, cudaMemcpyKind kind)
```

- Frees memory on the device.

```
__host__ __device__ cudaError_t cudaFree(void* devPtr)
```

- Allocates memory that will be automatically managed by the Unified Memory system.

```
__host__ cudaError_t cudaMallocManaged(void** devPtr, size_t size, unsigned int flags)
```



# CUDA Memory Management II

...

```
int* devArray;  
int size = 120;  
cudaMalloc(&devArray, size * sizeof(int));
```

```
// Send data to GPU  
cudaMemcpy(devArray, hostArrayA, size * sizeof(int),  
            cudaMemcpyHostToDevice);
```

```
// Process data  
kernel<<<...>>>(devArray, size, ...);
```

```
// Get result  
cudaMemcpy(hostArrayB, devArray, size * sizeof(int),  
            cudaMemcpyDeviceToHost);
```

```
cudaFree(devArray);
```

...

# CUDA Memory Management III

```
...  
  
int* unifiedArray;  
int size = 120;  
  
// Allocate on unified memory  
cudaMallocManaged(&unifiedArray, size * sizeof(int));  
  
// Process data  
kernel<<<...>>>(unifiedArray, size, ...);  
  
// Do something with the processed data  
host_func(unifiedArray, size);  
  
cudaFree(devArray);  
  
...
```

# CUDA Error Handling

- Every CUDA call (except kernel launches) return an error code of type `cudaError_t`

```
cudaError_t err = cudaMalloc(&fooPtr, -1);  
if (cudaSuccess != err)  
    printf("Error! -: -%d\n", err);
```

- CUDA kernel invocations do not return any value. Error from a CUDA kernel call can be checked after its execution by calling `cudaGetLastError()`

```
fooKernel<<<b,t>>>();  
cudaDeviceSynchronize(); // Remember to sync!  
cudaError_t err = cudaGetLastError();  
...
```

# CUDA Compilation

- Any source file containing CUDA language extensions must be compiled with **NVCC**
- NVCC is a **compiler driver**, It automatically invokes all the necessary tools and compilers like cudacc and g++
- Any executable with CUDA code requires the **CUDA runtime library** (cudart)
- Simple usage of NVCC to compile:

```
nvcc -std=c++11 <source> -o <executable>
```

Clone the repository at

`https://github.com/z1ko/parallel\_prog\_course`

Today you have the following exercises:

1. **hello\_world**: Simple program to test compilation.
2. **vector\_set**: Initialize a vector with a value.
3. **vector\_add**: Add two vectors.